

Recommendations_with_IBM

December 20, 2022

1 Recommendations with IBM

In this notebook, you will be putting your recommendation skills to use on real data from the IBM Watson Studio platform.

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project [RUBRIC](#). **Please save regularly.**

By following the table of contents, you will build out a number of different methods for making recommendations that can be used for different situations.

1.1 Table of Contents

I. Section ?? II. Section ?? III. Section ?? IV. Section ?? V. Section ?? VI. Section ??

At the end of the notebook, you will find directions for how to submit your work. Let's get started by importing the necessary libraries and reading in the data.

```
In [47]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import project_tests as t
import pickle

%matplotlib inline

df_user_interactions = pd.read_csv('data/user-item-interactions.csv')
df_articles = pd.read_csv('data/articles_community.csv')
del df_user_interactions['Unnamed: 0']
del df_articles['Unnamed: 0']

# Show df_user_interactions to get an idea of the data
print()
print("user-item-interactions - df_user_interactions.head() - ", df_user_interactions.h

user-item-interactions - df_user_interactions.head() -      article_id
0      1430.0  using pixiedust for fast, flexible, and easier...
1      1314.0      healthcare python streaming application demo
2      1429.0      use deep learning for image classification
```

```

3      1338.0          ml optimization using cognitive assistant
4      1276.0          deploy your python model as a restful api

```

```

                                email
0  ef5f11f77ba020cd36e1105a00ab868bbdbf7fe7
1  083cbdfa93c8444beaa4c5f5e0f5f9198e4f9e0b
2  b96a4f2e92d8572034b1e9b28f9ac673765cd074
3  06485706b34a5c9bf2a0ecdac41daf7e7654ceb7
4  f01220c46fc92c6e6b161b1849de11faacd7ccb2

```

```

In [48]: # Shape
print()
print("Shape of df_user_interactions.shape - ", df_user_interactions.shape)

# Columns
print()
print("Shape of df_user_interactions.columns - ", df_user_interactions.columns)

# Show df_articles to get an idea of the data
print()
print("df_articles.head() - df_articles() - ", df_articles.head())

# Shape
print()
print("Shape of df_articles.shape - ", df_articles.shape)

# Columns
print()
print("Shape of df_articles.columns - ", df_articles.columns)

print("\n\n --- Print few sample data --- ")

print()
print("df_user_interactions.iloc[0] = \n", df_user_interactions.iloc[0])

print()
print("df_user_interactions.iloc[1] = \n", df_user_interactions.iloc[1])

print()
print("df_articles.iloc[0] = \n", df_articles.iloc[0])

print()
print("df_articles.iloc[1] = \n", df_articles.iloc[1])

print()

```

```
print("df_articles.iloc[2] = \n", df_articles.iloc[2])
```

```
Shape of df_user_interactions.shape - (45993, 3)
```

```
Shape of df_user_interactions.columns - Index(['article_id', 'title', 'email'], dtype='object')
```

```
df_articles.head() - df_articles() - doc_body \
0 Skip navigation Sign in SearchLoading...\r\n\r...
1 No Free Hunch Navigation * kaggle.com\r\n\r\n ...
2 * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
3 DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...
4 Skip navigation Sign in SearchLoading...\r\n\r...
```

```
doc_description \
0 Detect bad readings in real time using Python ...
1 See the forest, see the trees. Here lies the c...
2 Heres this weeks news in Data Science and Bi...
3 Learn how distributed DBs solve the problem of...
4 This video demonstrates the power of IBM DataS...
```

	doc_full_name	doc_status	article_id
0	Detect Malfunctioning IoT Sensors with Streami...	Live	0
1	Communicating data science: A guide to present...	Live	1
2	This Week in Data Science (April 18, 2017)	Live	2
3	DataLayer Conference: Boost the performance of...	Live	3
4	Analyze NY Restaurant data using Spark in DSX	Live	4

```
Shape of df_articles.shape - (1056, 5)
```

```
Shape of df_articles.columns - Index(['doc_body', 'doc_description', 'doc_full_name', 'doc_stat',
    'article_id'],
    dtype='object')
```

```
--- Print few sample data ---
```

```
df_user_interactions.iloc[0] =
    article_id 1430
title using pixiedust for fast, flexible, and easier...
email ef5f11f77ba020cd36e1105a00ab868bbdbf7fe7
Name: 0, dtype: object
```

```
df_user_interactions.iloc[1] =
    article_id 1314
title healthcare python streaming application demo
email 083cbdfa93c8444beaa4c5f5e0f5f9198e4f9e0b
Name: 1, dtype: object
```

```
df_articles.iloc[0] =
  doc_body      Skip navigation Sign in SearchLoading...\r\n\r...
  doc_description Detect bad readings in real time using Python ...
  doc_full_name   Detect Malfunctioning IoT Sensors with Streami...
  doc_status      Live
  article_id      0
Name: 0, dtype: object
```

```
df_articles.iloc[1] =
  doc_body      No Free Hunch Navigation * kaggle.com\r\n\r\n ...
  doc_description See the forest, see the trees. Here lies the c...
  doc_full_name   Communicating data science: A guide to present...
  doc_status      Live
  article_id      1
Name: 1, dtype: object
```

```
df_articles.iloc[2] =
  doc_body      * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
  doc_description Heres this weeks news in Data Science and Bi...
  doc_full_name   This Week in Data Science (April 18, 2017)
  doc_status      Live
  article_id      2
Name: 2, dtype: object
```

1.1.1 Part I: Exploratory Data Analysis

Use the dictionary and cells below to provide some insight into the descriptive statistics of the data.

1. What is the distribution of how many articles a user interacts with in the dataset? Provide a visual and descriptive statistics to assist with giving a look at the number of times each user interacts with an article.

```
In [49]: # Total Interactions
total_interactions = df_user_interactions['email'].value_counts()

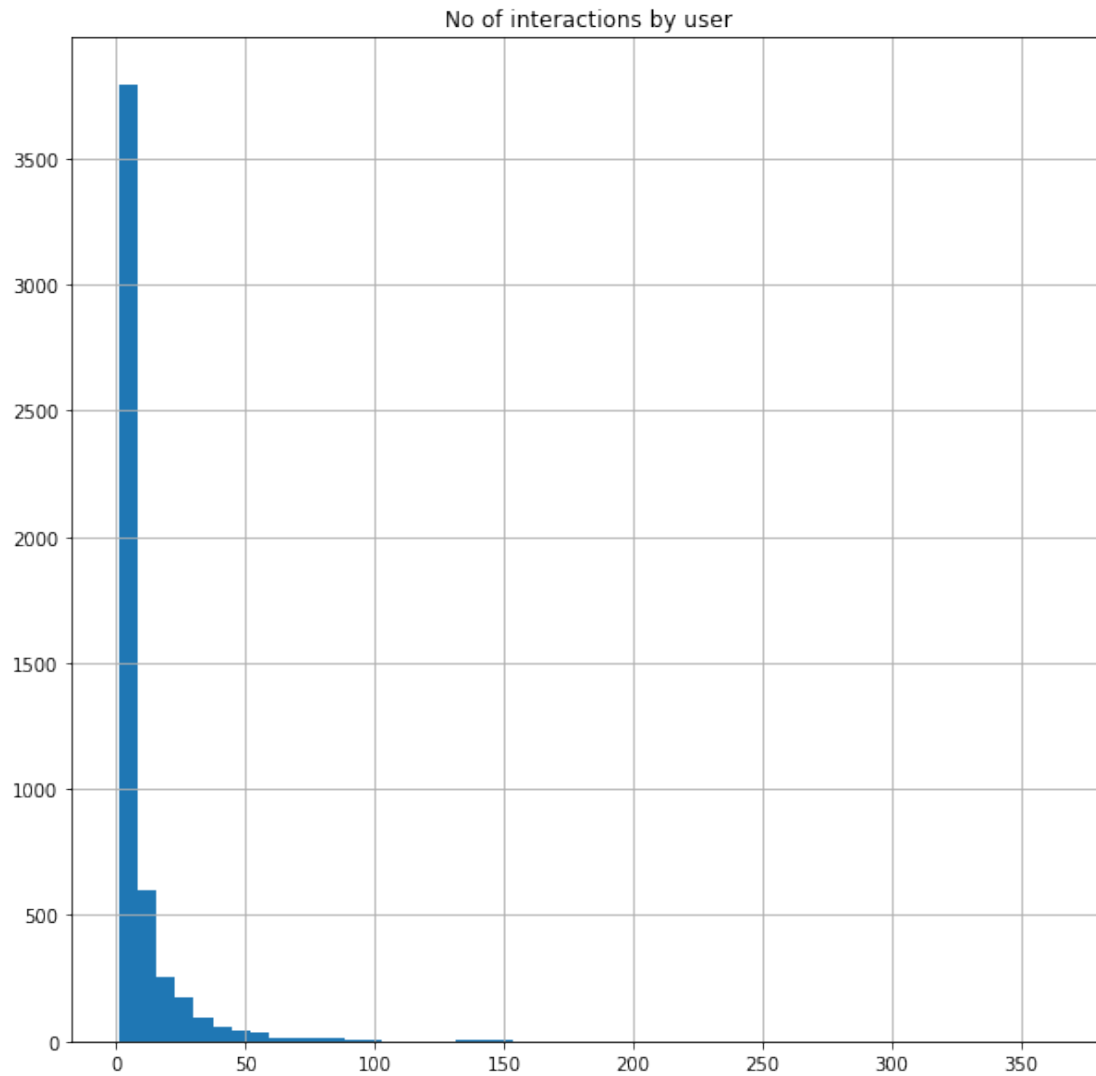
print()
print("Total Interactions Summary ==> \n", pd.DataFrame(total_interactions.describe()))

total_interactions.hist(bins = 50, grid = True, figsize = (10, 10))
plt.title('No of interactions by user')
```

```
Total Interactions Summary ==>
      email
count  5148.000000
mean    8.930847
```

```
std      16.802267
min       1.000000
25%       1.000000
50%       3.000000
75%       9.000000
max      364.000000
```

```
Out[49]: Text(0.5,1,'No of interactions by user')
```



```
In [50]: print("Most of the user have hardly any interaction with articles. May be they are new")
```

Most of the user have hardly any interaction with articles. May be they are new.

```
In [51]: # Fill in the median and maximum number of user_article interactions below
```

```
median_val = 3 # 50% of individuals interact with ____ number of articles or fewer.  
max_views_by_user = 364 # The maximum number of user-article interactions by any 1 user
```

2. Explore and remove duplicate articles from the **df_content** dataframe.

```
In [52]: # Check the shape of articles
```

```
print()  
print("Shape = ", df_articles.shape)
```

```
# Find and explore duplicate articles
```

```
print("Duplicated articles => \n", df_articles[df_articles['article_id'].duplicated()])
```

```
Shape = (1056, 5)
```

```
Duplicated articles =>
```

```
[ 50 221 398 577 232]
```

```
In [53]: # Remove any rows that have the same article_id - only keep the first
```

```
df_articles = df_articles[~df_articles['article_id'].duplicated(keep='first')]
```

```
# Check the shape of articles again
```

```
print("Shape after removing duplicates = ", df_articles.shape)
```

```
Shape after removing duplicates = (1051, 5)
```

3. Use the cells below to find:

a. The number of unique articles that have an interaction with a user.

b. The number of unique articles in the dataset (whether they have any interactions or not). c.

The number of unique users in the dataset. (excluding null values) d. The number of user-article interactions in the dataset.

```
In [54]: unique_articles = len(df_user_interactions['article_id'].unique()) # The number of unique
```

```
total_articles = len(df_articles) # The number of unique articles on the IBM platform
```

```
unique_users = len(df_user_interactions['email'].dropna().unique()) # df_user_interactions
```

```
user_article_interactions = len(df_user_interactions) # df_user_interactions.shape[0] #
```

4. Use the cells below to find the most viewed **article_id**, as well as how often it was viewed.

After talking to the company leaders, the email_mapper function was deemed a reasonable way to map users to ids. There were a small number of null values, and it was found that all of these null values likely belonged to a single user (which is how they are stored using the function below).

```
In [55]: most_viewed_article_id = str(df_user_interactions['article_id'].value_counts().idxmax())
```

```
max_views = df_user_interactions['article_id'].value_counts().max() # The most viewed a
```

```
In [56]: print("most_viewed_article_id ==> ", most_viewed_article_id)
```

```
print("max_views ==> ", max_views)
```

```
most_viewed_article_id ==> 1429.0
max_views ==> 937
```

```
In [57]: ## No need to change the code here - this will be helpful for later parts of the notebook
        # Run this cell to map the user email to a user_id column and remove the email column
```

```
def email_mapper():
    coded_dict = dict()
    cter = 1
    email_encoded = []

    for val in df_user_interactions['email']:
        if val not in coded_dict:
            coded_dict[val] = cter
            cter+=1

    email_encoded.append(coded_dict[val])
    return email_encoded

email_encoded = email_mapper()
del df_user_interactions['email']
df_user_interactions['user_id'] = email_encoded

# show header
print()
print("df_user_interactions.head() ==> \n", df_user_interactions.head())
```

```
df_user_interactions.head() ==>
```

	article_id	title	user_id
0	1430.0	using pixiedust for fast, flexible, and easier...	1
1	1314.0	healthcare python streaming application demo	2
2	1429.0	use deep learning for image classification	3
3	1338.0	ml optimization using cognitive assistant	4
4	1276.0	deploy your python model as a restful api	5

```
In [58]: ## If you stored all your results in the variable names above,
        ## you shouldn't need to change anything in this cell
```

```
sol_1_dict = {
    '50% of individuals have ____ or fewer interactions.': median_val,
    'The total number of user-article interactions in the dataset is ____.': user_a
    'The maximum number of user-article interactions by any 1 user is ____.': max_v
    'The most viewed article in the dataset was viewed ____ times.': max_views,
    'The article_id of the most viewed article is ____.': most_viewed_article_id,
    'The number of unique articles that have at least 1 rating ____.': unique_artic
```

```

        'The number of unique users in the dataset is _____': unique_users,
        'The number of unique articles on the IBM platform': total_articles
    }

    # Test your dictionary against the solution
    t.sol_1_test(sol_1_dict)

```

It looks like you have everything right here! Nice job!

1.1.2 Part II: Rank-Based Recommendations

Unlike in the earlier lessons, we don't actually have ratings for whether a user liked an article or not. We only know that a user has interacted with an article. In these cases, the popularity of an article can really only be based on how often an article was interacted with.

1. Fill in the function below to return the `n` top articles ordered with most interactions as the top. Test your function using the tests below.

```

In [59]: def get_top_articles(n, df=df_user_interactions):
    """
    INPUT:
    n - (int) the number of top articles to return
    df - (pandas dataframe) df as defined at the top of the notebook

    OUTPUT:
    top_articles - (list) A list of the top 'n' article titles

    """
    # Your code here
    top_articles = list(df_user_interactions['title'].value_counts().reset_index().head(n))

    return top_articles # Return the top article titles from df (not df_content)

def get_top_article_ids(n, df=df_user_interactions):
    """
    INPUT:
    n - (int) the number of top articles to return
    df - (pandas dataframe) df as defined at the top of the notebook

    OUTPUT:
    top_articles - (list) A list of the top 'n' article titles

    """
    # Your code here
    top_articles = list(df_user_interactions['article_id'].value_counts().reset_index().head(n))

    return top_articles # Return the top article ids

```



```
In [60]: print("get_top_articles(10) ==> \n", get_top_articles(10))
         print()
         print("get_top_article_ids(10) ==> \n", get_top_article_ids(10))

get_top_articles(10) ==>
['use deep learning for image classification', 'insights from new york car accident reports', '

get_top_article_ids(10) ==>
[1429.0, 1330.0, 1431.0, 1427.0, 1364.0, 1314.0, 1293.0, 1170.0, 1162.0, 1304.0]
```

```
In [61]: # Test your function by returning the top 5, 10, and 20 articles
         top_5 = get_top_articles(5)
         top_10 = get_top_articles(10)
         top_20 = get_top_articles(20)

         # Test each of your three lists from above
         t.sol_2_test(get_top_articles)
```

Your top_5 looks like the solution list! Nice job.
 Your top_10 looks like the solution list! Nice job.
 Your top_20 looks like the solution list! Nice job.

1.1.3 Part III: User-User Based Collaborative Filtering

1. Use the function below to reformat the **df** dataframe to be shaped with users as the rows and articles as the columns.

- Each **user** should only appear in each **row** once.
- Each **article** should only show up in one **column**.
- If a user has interacted with an article, then place a 1 where the user-row meets for that article-column. It does not matter how many times a user has interacted with the article, all entries where a user has interacted with an article should be a 1.
- If a user has not interacted with an item, then place a zero where the user-row meets for that article-column.

Use the tests to make sure the basic structure of your matrix matches what is expected by the solution.

```
In [62]: # create the user-article matrix with 1's and 0's

def create_user_item_matrix(df):
    """
    INPUT:
    df - pandas dataframe with article_id, title, user_id columns
```

OUTPUT:

user_item - user item matrix

Description:

Return a matrix with user ids as rows and article ids on the columns with 1 values an article and a 0 otherwise

'''

Fill in the function here

`df['article_id'] = df['article_id'].astype(str)`

`user_item_pivot = df.groupby(['user_id', 'article_id'])['title'].count().notnull().`

`user_item_matrix = user_item_pivot.notnull().astype(np.int64)`

`return user_item_matrix # return the user_item matrix`

`user_item_matrix = create_user_item_matrix(df_user_interactions)`

In [63]: *## Tests: You should just need to run this cell. Don't change the code.*

`assert user_item_matrix.shape[0] == 5149, "Oops! The number of users in the user-artic`

`assert user_item_matrix.shape[1] == 714, "Oops! The number of articles in the user-art`

`assert user_item_matrix.sum(axis=1)[1] == 36, "Oops! The number of articles seen by us`

`print("You have passed our quick tests! Please proceed!")`

You have passed our quick tests! Please proceed!

2. Complete the function below which should take a `user_id` and provide an ordered list of the most similar users to that user (from most similar to least similar). The returned result should not contain the provided `user_id`, as we know that each user is similar to him/herself. Because the results for each user here are binary, it (perhaps) makes sense to compute similarity as the dot product of two users.

Use the tests to test your function.

In [64]: `def find_similar_users(user_id, user_item=user_item_matrix):`

'''

INPUT:

user_id - (int) a user_id

user_item - (pandas dataframe) matrix of users by articles:

1's when a user has interacted with an article, 0 otherwise

OUTPUT:

similar_users - (list) an ordered list where the closest users (largest dot product) are listed first

Description:

Computes the similarity of every pair of users based on the dot product

Returns an ordered

'''

```

no_users = user_item_matrix.shape[0]
user_row = user_item_matrix.loc[user_id]
user_tile = np.tile(user_row.values, (no_users, 1))

# compute similarity of each user to the provided user
similarities = np.multiply(user_tile, user_item_matrix).sum(axis = 1)

# sort by similarity
similarities.sort_values(ascending=False, inplace = True)

# create list of just the ids
most_similar_users = list(similarities.index)

# remove the own user's id
most_similar_users.remove(user_id)

return most_similar_users # return a list of the users in order from most to least

```

In [65]: # Do a spot check of your function

```

print("The 10 most similar users to user 1 are: {}".format(find_similar_users(1)[:10]))
print("The 5 most similar users to user 3933 are: {}".format(find_similar_users(3933)[:5]))
print("The 3 most similar users to user 46 are: {}".format(find_similar_users(46)[:3]))

```

The 10 most similar users to user 1 are: [3933, 23, 3782, 203, 4459, 131, 3870, 46, 4201, 5041]
The 5 most similar users to user 3933 are: [1, 23, 3782, 4459, 203]
The 3 most similar users to user 46 are: [4201, 23, 3782]

3. Now that you have a function that provides the most similar users to each user, you will want to use these users to find articles you can recommend. Complete the functions below to return the articles you would recommend to each user.

In [66]: def get_article_names(article_ids, df=df_user_interactions):

'''

INPUT:

article_ids - (list) a list of article ids

df - (pandas dataframe) df as defined at the top of the notebook

OUTPUT:

article_names - (list) a list of article names associated with the list of article ids
(this is identified by the title column)

'''

Your code here

```

article_names = (df_user_interactions[df_user_interactions['article_id'].isin(article_ids)]
                 .drop_duplicates().values.tolist())

```

```

return article_names # Return the article names associated with list of article ids

```

```

print("get_article_names([1430.0, 1314.0]) ==> \n", get_article_names(['1430.0', '1314.0']))

def get_user_articles(user_id, user_item=user_item_matrix):
    """
    INPUT:
    user_id - (int) a user id
    user_item - (pandas dataframe) matrix of users by articles:
                1's when a user has interacted with an article, 0 otherwise

    OUTPUT:
    article_ids - (list) a list of the article ids seen by the user
    article_names - (list) a list of article names associated with the list of article
                    (this is identified by the doc_full_name column in df_content)

    Description:
    Provides a list of the article_ids and article titles that have been seen by a user
    """
    # Your code here
    # Find all the user ids that a certain user has seen
    user_vector = user_item_matrix.loc[user_id]
    user_vector_seen = user_vector.where(user_vector == 1).dropna()
    article_ids = list(user_vector_seen.index)

    # Find all the article names based on the article_ids returned in the last step
    article_names = get_article_names(article_ids)

    return article_ids, article_names # return the ids and names

print()
print("get_user_articles(1) ==> \n", get_user_articles(1))

import itertools

def user_user_recs(user_id, m=10):
    """
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them as recs
    Does this until m recommendations are found
    """

```

Notes:

Users who are the same closeness are chosen arbitrarily as the 'next' user

For the user where the number of recommended articles starts below m and ends exceeding m, the last items are chosen arbitrarily

```
'''
# Your code here
# Find all articles that the user has read before
user_read_articles = get_user_articles(user_id)[0]
similar_users = find_similar_users(user_id)

# Find all articles that has been read by the similiar users
read_articles = [get_user_articles(user)[0] for user in similar_users]
read_articles_list = list(itertools.chain.from_iterable(read_articles))

# remove duplicated read articles
read_articles_unique = pd.Series(read_articles_list).drop_duplicates().tolist()

# remove articles that has been seen by the given user, use the remaining articles
recs = [i for i in read_articles_unique if i not in user_read_articles][:m]

return recs # return your recommendations for this user_id

get_article_names([1430.0, 1314.0]) ==>
['using pixiedust for fast, flexible, and easier data analysis and experimentation', 'healthcar

get_user_articles(1) ==>
(['1052.0', '109.0', '1170.0', '1183.0', '1185.0', '1232.0', '1293.0', '1305.0', '1363.0', '136

In [67]: # Check Results
get_article_names(user_user_recs(1, 10)) # Return 10 recommendations for user 1
print()
print("---> Done with get_article_names()...")

---> Done with get_article_names()...

In [68]: # Test your functions here - No need to change this code - just run this cell
assert set(get_article_names(['1024.0', '1176.0', '1305.0', '1314.0', '1422.0', '1427.0
assert set(get_article_names(['1320.0', '232.0', '844.0'])) == set(['housing (2015): un
assert set(get_user_articles(20)[0]) == set(['1320.0', '232.0', '844.0'])
assert set(get_user_articles(20)[1]) == set(['housing (2015): united states demographic
assert set(get_user_articles(2)[0]) == set(['1024.0', '1176.0', '1305.0', '1314.0', '14
assert set(get_user_articles(2)[1]) == set(['using deep learning to reconstruct high-re
print("If this is all you see, you passed all of our tests! Nice job!")
```

If this is all you see, you passed all of our tests! Nice job!

4. Now we are going to improve the consistency of the **user_user_recs** function from above.

- Instead of arbitrarily choosing when we obtain users who are all the same closeness to a given user - choose the users that have the most total article interactions before choosing those with fewer article interactions.
- Instead of arbitrarily choosing articles from the user where the number of recommended articles starts below m and ends exceeding m, choose articles with the articles with the most total interactions before choosing those with fewer total interactions. This ranking should be what would be obtained from the **top_articles** function you wrote earlier.

```
In [70]: def get_top_sorted_users(user_id, df=df_user_interactions, user_item=user_item_matrix):  
    '''  
    INPUT:  
    user_id - (int)  
    df - (pandas dataframe) df as defined at the top of the notebook  
    user_item - (pandas dataframe) matrix of users by articles:  
                1's when a user has interacted with an article, 0 otherwise  
  
    OUTPUT:  
    neighbors_df - (pandas dataframe) a dataframe with:  
                    neighbor_id - is a neighbor user_id  
                    similarity - measure of the similarity of each user to the provided  
                    num_interactions - the number of articles viewed by the user - if a  
  
    Other Details - sort the neighbors_df by the similarity and then by number of inter  
                    highest of each is higher in the dataframe  
  
    '''  
    # Your code here  
    no_users = user_item_matrix.shape[0]  
    user_row = user_item_matrix.loc[user_id]  
    user_tile = np.tile(user_row.values, (no_users, 1))  
  
    # compute similarity  
    similarities = pd.DataFrame(np.multiply(user_tile, user_item_matrix).sum(axis = 1),  
                                columns = ['similarity'])  
  
    # Add column for total interactions per user  
    similarities['no_interactions'] = df_user_interactions.groupby('user_id').count()['  
  
    # sort on similarities, total number of interactions  
    similarities.sort_values(by = ['similarity', 'no_interactions'],  
                             ascending=False, inplace = True)
```

```

    # remove the user which is in input
    similarities.drop(user_id, inplace = True)

    # create neighbors_df
    neighbors_df = similarities.reset_index()
    neighbors_df.columns = ['neighbor_id', 'similarity', 'no_interactions']

    return neighbors_df # Return the dataframe specified in the doc_string

def user_user_recs_part2(user_id, m=10):
    """
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user by article id
    rec_names - (list) a list of recommendations for the user by article title

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them as recs
    Does this until m recommendations are found

    Notes:
    * Choose the users that have the most total article interactions
    before choosing those with fewer article interactions.

    * Choose articles with the articles with the most total interactions
    before choosing those with fewer total interactions.

    """
    # Your code here

    user_articles, _ = get_user_articles(user_id)

    neighbors_df = get_top_sorted_users(user_id)

    recs = []
    for each_user in neighbors_df['neighbor_id'].tolist():
        if len(recs) < m:
            possible_recs, _ = get_user_articles(each_user)

            for each_rec in possible_recs:

                if len(recs) < m:
                    if each_rec not in user_articles:

```

```

        if each_rec not in recs:
            recs.append(each_rec)
        else:
            break;
    else:
        break;

    rec_names = get_article_names(recs)

    return recs, rec_names

```

```

In [71]: # Quick spot check - don't change this code - just use it to test your functions
rec_ids, rec_names = user_user_recs_part2(20, 10)
print("The top 10 recommendations for user 20 are the following article ids:")
print(rec_ids)
print()
print("The top 10 recommendations for user 20 are the following article names:")
print(rec_names)

```

The top 10 recommendations for user 20 are the following article ids:

```
['1024.0', '1085.0', '109.0', '1150.0', '1151.0', '1152.0', '1153.0', '1154.0', '1157.0', '1160.0']
```

The top 10 recommendations for user 20 are the following article names:

```
['airbnb data for analytics: washington d.c. listings', 'analyze accident reports on amazon emr']
```

5. Use your functions from above to correctly fill in the solutions to the dictionary below. Then test your dictionary against the solution. Provide the code you need to answer each following the comments below.

```

In [72]: ### Tests with a dictionary of results

```

```

# Find the user that is most similar to user 1
user1_most_sim = get_top_sorted_users(1)['neighbor_id'].head(1).values[0]
print()
print("user1_most_sim ==> ", user1_most_sim)
# Find the 10th most similar user to user 131
user131_10th_sim = get_top_sorted_users(131)['neighbor_id'].head(10).values[-1]
print()
print("user131_10th_sim ==> ", user131_10th_sim)

```

```
user1_most_sim ==> 3933
```

```
user131_10th_sim ==> 242
```

```

In [73]: ## Dictionary Test Here
sol_5_dict = {

```



```

        'The user that is most similar to user 1.': user1_most_sim,
        'The user that is the 10th most similar to user 131': user131_10th_sim,
    }
    print("\n\n")
    print("TYPE ==> ", type(sol_5_dict))
    print("\n sol_5_dict.items() ==> \n", sol_5_dict.items())

    print("sol_5_dict['The user that is most similar to user 1.'] = ", sol_5_dict['The user
    print("sol_5_dict['The user that is the 10th most similar to user 131'] = ", sol_5_dict

    t.sol_5_test(sol_5_dict)

```

```
TYPE ==> <class 'dict'>
```

```

sol_5_dict.items() ==>
dict_items([('The user that is most similar to user 1.', 3933), ('The user that is the 10th mos
sol_5_dict['The user that is most similar to user 1.'] = 3933
sol_5_dict['The user that is the 10th most similar to user 131'] = 242
This all looks good! Nice job!

```

6. If we were given a new user, which of the above functions would you be able to use to make recommendations? Explain. Can you think of a better way we might make recommendations? Use the cell below to explain a better method for new users.

```

In [74]: '''
        For a new user, it is better to use the function get_top_articles(), because it is the
        '''

```

```
Out[74]: '\nFor a new user, it is better to use the function get_top_articles(), because it is t
```

Provide your response here.

7. Using your existing functions, provide the top 10 recommended articles you would provide for the a new user below. You can test your function against our thoughts to make sure we are all on the same page with how we might make a recommendation.

```

In [75]: new_user = '0.0'

        # What would your recommendations be for this new user '0.0'? As a new user, they have
        # Provide a list of the top 10 article ids you would give to
        new_user_recs = get_top_article_ids(10) # Your recommendations here

In [76]: assert set(new_user_recs) == set(['1314.0', '1429.0', '1293.0', '1427.0', '1162.0', '1364.0'

        print("That's right! Nice job!")

```

That's right! Nice job!

1.1.4 Part IV: Content Based Recommendations (EXTRA - NOT REQUIRED)

Another method we might use to make recommendations is to perform a ranking of the highest ranked articles associated with some term. You might consider content to be the **doc_body**, **doc_description**, or **doc_full_name**. There isn't one way to create a content based recommendation, especially considering that each of these columns hold content related information.

1. Use the function body below to create a content based recommender. Since there isn't one right answer for this recommendation tactic, no test functions are provided. Feel free to change the function inputs if you decide you want to try a method that requires more input values. The input values are currently set with one idea in mind that you may use to make content based recommendations. One additional idea is that you might want to choose the most popular recommendations that meet your 'content criteria', but again, there is a lot of flexibility in how you might make these recommendations.

1.1.5 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [ ]: def make_content_recs():  
        '''  
        INPUT:  
  
        OUTPUT:  
  
        '''
```

2. Now that you have put together your content-based recommendation system, use the cell below to write a summary explaining how your content based recommender works. Do you see any possible improvements that could be made to your function? Is there anything novel about your content based recommender?

1.1.6 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

Write an explanation of your content based recommendation system here.

3. Use your content-recommendation system to make recommendations for the below scenarios based on the comments. Again no tests are provided here, because there isn't one right answer that could be used to find these content based recommendations.

1.1.7 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [ ]: # make recommendations for a brand new user  
  
        # make a recommendations for a user who only has interacted with article id '1427.0'
```

1.1.8 Part V: Matrix Factorization

In this part of the notebook, you will build use matrix factorization to make article recommendations to the users on the IBM Watson Studio platform.

1. You should have already created a **user_item** matrix above in **question 1** of **Part III** above. This first question here will just require that you run the cells to get things set up for the rest of **Part V** of the notebook.

```
In [77]: # Load the matrix here
```

```
user_item_matrix = pd.read_pickle('user_item_matrix.p')
```

```
In [78]: # quick look at the matrix
```

```
print("user_item_matrix.head() ==> \n", user_item_matrix.head())
```

```
user_item_matrix.head() ==>
```

```
  article_id  0.0  100.0  1000.0  1004.0  1006.0  1008.0  101.0  1014.0  1015.0  \
user_id
1          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
2          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
3          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
4          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
5          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
```

```
  article_id  1016.0  ...    977.0  98.0  981.0  984.0  985.0  986.0  990.0  \
user_id          ...
1          0.0  ...    0.0  0.0  1.0  0.0  0.0  0.0  0.0
2          0.0  ...    0.0  0.0  0.0  0.0  0.0  0.0  0.0
3          0.0  ...    1.0  0.0  0.0  0.0  0.0  0.0  0.0
4          0.0  ...    0.0  0.0  0.0  0.0  0.0  0.0  0.0
5          0.0  ...    0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
  article_id  993.0  996.0  997.0
user_id
1          0.0    0.0    0.0
2          0.0    0.0    0.0
3          0.0    0.0    0.0
4          0.0    0.0    0.0
5          0.0    0.0    0.0
```

```
[5 rows x 714 columns]
```

2. In this situation, you can use Singular Value Decomposition from [numpy](#) on the user-item matrix. Use the cell to perform SVD, and explain why this is different than in the lesson.

```
In [79]: # Perform SVD on the User-Item Matrix Here
```

```
u, s, vt = np.linalg.svd(user_item_matrix) # use the built in to get the three matrices
```

```
In [80]: ''' If there were missing data, doing SVD will lead to errors. But here we do not have
```

```
Out[80]: ' If there were missing data, doing SVD will lead to errors. But here we do not have th
```

Provide your response here.

3. Now for the tricky part, how do we choose the number of latent features to use? Running the below cell, you can see that as the number of latent features increases, we obtain a lower error rate on making predictions for the 1 and 0 values in the user-item matrix. Run the cell below to get an idea of how the accuracy improves as we increase the number of latent features.

```
In [81]: num_latent_feats = np.arange(10,700+10,20)
        sum_errs = []


        for k in num_latent_feats:
            # restructure with k latent features
            s_new, u_new, vt_new = np.diag(s[:k]), u[:, :k], vt[:k, :]

            # take dot product
            user_item_est = np.around(np.dot(np.dot(u_new, s_new), vt_new))

            # compute error for each prediction to actual value
            diffs = np.subtract(user_item_matrix, user_item_est)

            # total errors and keep track of them
            err = np.sum(np.sum(np.abs(diffs)))
            sum_errs.append(err)

        plt.plot(num_latent_feats, 1 - np.array(sum_errs)/df_user_interactions.shape[0]);
        plt.xlabel('Number of Latent Features');
        plt.ylabel('Accuracy');
        plt.title('Accuracy vs. Number of Latent Features');
```



4. From the above, we can't really be sure how many features to use, because simply having a better way to predict the 1's and 0's of the matrix doesn't exactly give us an indication of if we are able to make good recommendations. Instead, we might split our dataset into a training and test set of data, as shown in the cell below.

Use the code from question 3 to understand the impact on accuracy of the training and test sets of data with different numbers of latent features. Using the split below:

- How many users can we make predictions for in the test set?
- How many users are we not able to make predictions for because of the cold start problem?
- How many articles can we make predictions for in the test set?
- How many articles are we not able to make predictions for because of the cold start problem?

```
In [82]: df_train = df_user_interactions.head(40000)
         df_test = df_user_interactions.tail(5993)

def create_test_and_train_user_item(df_train, df_test):
    """
    INPUT:
    df_train - training dataframe
    df_test - test dataframe

    OUTPUT:
    user_item_train - a user-item matrix of the training dataframe
```

```

        (unique users for each row and unique articles for each column)
user_item_test - a user-item matrix of the testing dataframe
        (unique users for each row and unique articles for each column)
test_idx - all of the test user ids
test_arts - all of the test article ids

'''
# Your code here
user_item_train = create_user_item_matrix(df_train)
user_item_test = create_user_item_matrix(df_test)

test_idx = list(user_item_test.index)
test_arts = list(user_item_test.columns)

return user_item_train, user_item_test, test_idx, test_arts

```

```

user_item_train, user_item_test, test_idx, test_arts = create_test_and_train_user_item(

```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:16: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

```

app.launch_new_instance()

```

In [83]: # Replace the values in the dictionary below

```

a = 662
b = 574
c = 20
d = 0

```

```

sol_4_dict = {
    'How many users can we make predictions for in the test set?': c,
    'How many users in the test set are we not able to make predictions for because of': b,
    'How many articles can we make predictions for in the test set?': b,
    'How many articles in the test set are we not able to make predictions for because of': c,
}

t.sol_4_test(sol_4_dict)

```

Awesome job! That's right! All of the test articles are in the training data, but there are on

5. Now use the **user_item_train** dataset from above to find U, S, and V transpose using SVD. Then find the subset of rows in the **user_item_test** dataset that you can predict using this matrix decomposition with different numbers of latent features to see how many features makes sense

to keep based on the accuracy on the test data. This will require combining what was done in questions 2 - 4.

Use the cells below to explore how well SVD works towards making predictions for recommendations on the test data.

```
In [84]: # fit SVD on the user_item_train matrix
         u_train, s_train, vt_train = np.linalg.svd(user_item_train) # fit svd similar to above

In [85]: # Use these cells to see how well you can use the training
         # decomposition to predict on test data

         test_rows_idx = user_item_train.index.isin(test_idx)
         test_col_idxes = user_item_train.columns.isin(test_arts)
         u_test = u_train[test_rows_idx, :]
         vt_test = vt_train[:, test_col_idxes]
         # find the users that exists in both training and test datasets
         user_present_both = np.intersect1d(user_item_test.index, user_item_train.index)
         user_item_test_predictable = user_item_test[user_item_test.index.isin(user_present_both)]

         # initialize testing parameters
         num_latent_feats = np.arange(10, 700+10, 20)
         sum_errs_train = []
         sum_errs_test = []

         for k in num_latent_feats:
             # restructure with k latent features for both training and test sets
             s_train_lat, u_train_lat, vt_train_lat = np.diag(s_train[:k]), u_train[:, :k], vt_train[:, :k]
             u_test_lat, vt_test_lat = u_test[:, :k], vt_test[:, :k]

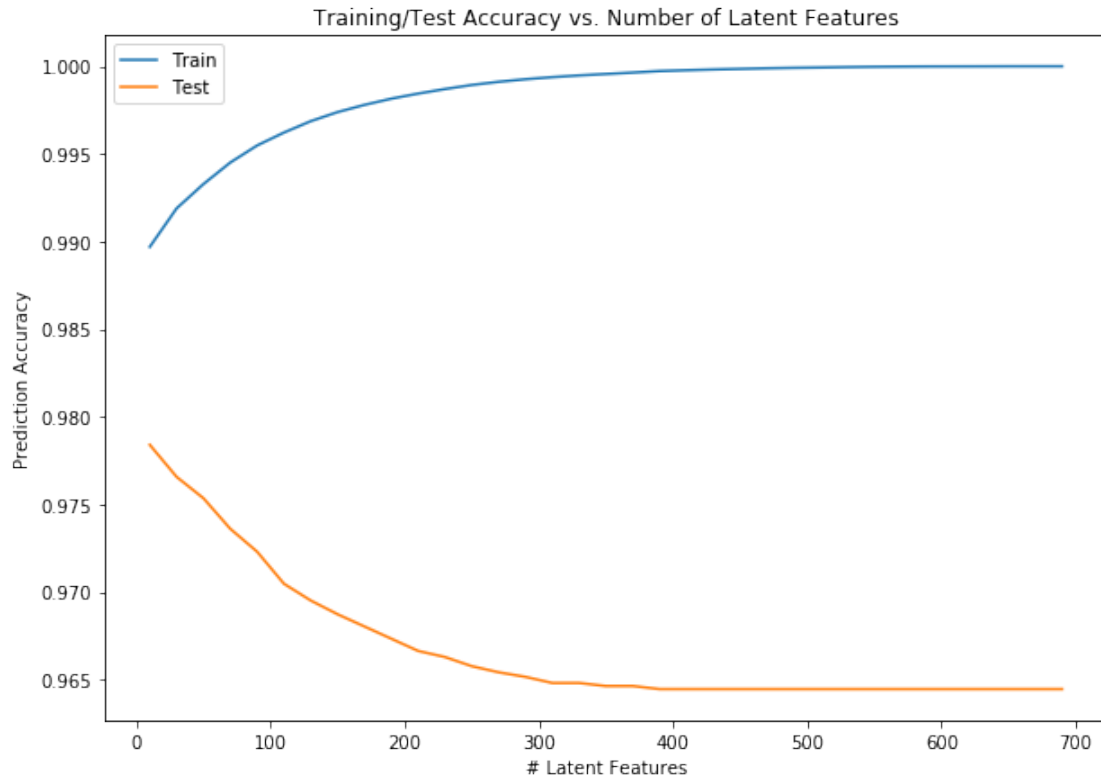
             # take dot product for both training and test sets
             user_item_train_est = np.around(np.dot(np.dot(u_train_lat, s_train_lat), vt_train_lat))
             user_item_test_est = np.around(np.dot(np.dot(u_test_lat, s_train_lat), vt_test_lat))

             # compute error for each prediction to actual value
             diffs_train = np.subtract(user_item_train, user_item_train_est)
             diffs_test = np.subtract(user_item_test_predictable, user_item_test_est)

             # total errors and keep track of them for both training and test sets
             err_train = np.sum(np.sum(np.abs(diffs_train)))
             err_test = np.sum(np.sum(np.abs(diffs_test)))
             sum_errs_train.append(err_train)
             sum_errs_test.append(err_test)

         plt.figure(figsize=(10, 7))
         plt.plot(num_latent_feats, 1 - np.array(sum_errs_train)/(user_item_train.shape[0]
                                                                    * user_item_test_predictable.shape[0]))
         plt.plot(num_latent_feats, 1 - np.array(sum_errs_test)/(user_item_test_predictable.shape[0]
                                                                    * user_item_test_predictable.shape[0]))
         plt.xlabel('# Latent Features')
```

```
plt.ylabel('Prediction Accuracy')
plt.legend()
plt.title('Training/Test Accuracy vs. Number of Latent Features')
plt.show()
```



6. Use the cell below to comment on the results you found in the previous question. Given the circumstances of your results, discuss what you might do to determine if the recommendations you make with any of the above recommendation systems are an improvement to how users currently find articles?

Your response here.

In [86]: '''

1. There are very less number of cases we have in the test set (20). This will create i
have to use rank based recommendation. We should increase training data accuracy and de
increasing latent features.

2. For the second part we should try A/B testing technique. We might try segregating us
different version of the site with some recommendation functionality and suggesting the
this for a while we can compare various metrics like interaction numbers, number of pag
recommendations are similar or diffierent.

'''

Out[86]: '\n1. There are very less number of cases we have in the test set (20). This will creat

Extras Using your workbook, you could now save your recommendations for each user, develop a class to make new predictions and update your results, and make a flask app to deploy your results. These tasks are beyond what is required for this project. However, from what you learned in the lessons, you are certainly capable of taking these tasks on to improve upon your work here!

1.2 Conclusion

Congratulations! You have reached the end of the Recommendations with IBM project!

Tip: Once you are satisfied with your work here, check over your report to make sure that it satisfies all the areas of the [rubric](#). You should also probably remove all of the "Tips" like this one so that the presentation is as polished as possible.

1.3 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File > Download as** sub-menu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```
In [87]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Recommendations_with_IBM.ipynb'])
```

```
Out[87]: 0
```

```
In [ ]:
```